

In the course you learned how to do classification using Fashion MNIST, a data set containing items of clothing. There's another, similar dataset called MNIST which has items of handwriting -- the digits 0 through 9.

Write an MNIST classifier that trains to 99% accuracy or above, and does it without a fixed number of epochs -- i.e. you should stop training once you reach that level of accuracy.

Some notes:

- 1. It should succeed in less than 10 epochs, so it is okay to change epochs to 10, but nothing larger
- 2. When it reaches 99% or greater it should print out the string "Reached 99% accuracy so cancelling training!"
- 3. If you add any additional variables, make sure you use the same names as the ones used in the class

I've started the code for you below -- how would you finish it?

Things to Note:

- 1. When coding the `class myCallback`, Python 3 will run into an error

```
TypeError: '>' not supported between instances of 'NoneType' and 'float'
```

when using the code

```
if(logs.get('accuracy')>0.99):
```

For Python 3, use the following equivalent code line

```
if logs.get('accuracy') is not None and logs.get('accuracy') > 0.99:
```

- 2. You can run the notebook using TensorFlow 2.5.0

```
import tensorflow as tf
```

```
print(tf.__version__)
```

2.6.0

```
# mnist = tf.keras.datasets.mnist
# (x_train, y_train),(x_test, y_test) = mnist.load_data()
# import numpy as np
# np.set_printoptions(linewidth=200)
# print(x_train[4])
# print(y_train[4])
# import matplotlib.pyplot as plt
# plt.imshow(x_train[4])
```

```
# GRADED FUNCTION: train_mnist
```

```
def train_mnist():
```

```
    class myCallback(tf.keras.callbacks.Callback):
        def on_epoch_end(self, epoch, logs={}):
            if logs.get('accuracy') is not None and logs.get('accuracy') > 0.99:
                print("\nReached 99% accuracy so cancelling training!")
                self.model.stop_training = True
```

```
    mnist = tf.keras.datasets.mnist
```

```
    (x_train, y_train),(x_test, y_test) = mnist.load_data()
```

```
    mnist = tf.keras.datasets.mnist
```

```
    (x_train, y_train),(x_test, y_test) = mnist.load_data()
    x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
    callbacks = myCallback()
```

```
    model = tf.keras.models.Sequential([
        tf.keras.layers.Flatten(input_shape=(28, 28)), # Adding a layer with single channel image 28x28
        tf.keras.layers.Dense(512, activation=tf.nn.relu), # Adds a layer of neurons
        tf.keras.layers.Dense(10, activation=tf.nn.softmax) # Takes a set of values and effectively picks the biggest one between 0 or 1
    ])
```

```
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
# model fitting
history = model.fit(x_train, y_train, epochs=10, callbacks=[callbacks] # CallBack function check and stop epoch once the model accuracy exceeding
)
# model fitting

return history.epoch, history.history['accuracy'][-1]
```

```
history.epoch, history.history['accuracy'][-1]train_mnist()
```

```
↳ Epoch 1/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.2003 - accuracy: 0.9414
Epoch 2/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0788 - accuracy: 0.9762
Epoch 3/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0516 - accuracy: 0.9836
Epoch 4/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0359 - accuracy: 0.9880
Epoch 5/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0265 - accuracy: 0.9913

Reached 99% accuracy so cancelling training!
([0, 1, 2, 3, 4], 0.9912833571434021)
```

