

```
# Import all the necessary files!
import os
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras import Model

# Download the inception v3 weights
!wget --no-check-certificate \
    https://storage.googleapis.com/mledu-datasets/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5 \
    -O /tmp/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5

# Import the inception model
from tensorflow.keras.applications.inception_v3 import InceptionV3

# Create an instance of the inception model from the local pre-trained weights
local_weights_file = '/tmp/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5'

pre_trained_model = InceptionV3(input_shape = (150, 150, 3), ### YOUR CODE HERE
                                include_top = False,
                                weights = None)

pre_trained_model.load_weights(local_weights_file)

# Make all the layers in the pre-trained model non-trainable
for layer in pre_trained_model.layers:
    layer.trainable = False ### YOUR CODE HERE

# Print the model summary
pre_trained_model.summary()
```

batch_normalization_184 (BatchN	(None, 3, 3, 384)	1152	conv2d_184[0][0]
activation_180 (Activation)	(None, 3, 3, 384)	0	batch_normalization_180[0][0]
activation_184 (Activation)	(None, 3, 3, 384)	0	batch_normalization_184[0][0]
conv2d_181 (Conv2D)	(None, 3, 3, 384)	442368	activation_180[0][0]



conv2d_182 (Conv2D)	(None, 3, 3, 384)	442368	activation_180[0][0]
conv2d_185 (Conv2D)	(None, 3, 3, 384)	442368	activation_184[0][0]
conv2d_186 (Conv2D)	(None, 3, 3, 384)	442368	activation_184[0][0]
average_pooling2d_17 (AveragePo	(None, 3, 3, 2048)	0	mixed9[0][0]
conv2d_179 (Conv2D)	(None, 3, 3, 320)	655360	mixed9[0][0]
batch_normalization_181 (BatchN	(None, 3, 3, 384)	1152	conv2d_181[0][0]
batch_normalization_182 (BatchN	(None, 3, 3, 384)	1152	conv2d_182[0][0]
batch_normalization_185 (BatchN	(None, 3, 3, 384)	1152	conv2d_185[0][0]
batch_normalization_186 (BatchN	(None, 3, 3, 384)	1152	conv2d_186[0][0]
conv2d_187 (Conv2D)	(None, 3, 3, 192)	393216	average_pooling2d_17[0][0]
batch_normalization_179 (BatchN	(None, 3, 3, 320)	960	conv2d_179[0][0]
activation_181 (Activation)	(None, 3, 3, 384)	0	batch_normalization_181[0][0]
activation_182 (Activation)	(None, 3, 3, 384)	0	batch_normalization_182[0][0]
activation_185 (Activation)	(None, 3, 3, 384)	0	batch_normalization_185[0][0]
activation_186 (Activation)	(None, 3, 3, 384)	0	batch_normalization_186[0][0]
batch_normalization_187 (BatchN	(None, 3, 3, 192)	576	conv2d_187[0][0]
activation_179 (Activation)	(None, 3, 3, 320)	0	batch_normalization_179[0][0]
mixed9_1 (Concatenate)	(None, 3, 3, 768)	0	activation_181[0][0] activation_182[0][0]
concatenate_3 (Concatenate)	(None, 3, 3, 768)	0	activation_185[0][0] activation_186[0][0]
activation_187 (Activation)	(None, 3, 3, 192)	0	batch_normalization_187[0][0]
mixed10 (Concatenate)	(None, 3, 3, 2048)	0	activation_179[0][0] mixed9_1[0][0] concatenate_3[0][0] activation_187[0][0]

=====
Total params: 21,802,784

Trainable params: 0
Non-trainable params: 21,802,784

```
last_layer = pre_trained_model.get_layer('mixed7') ### YOUR CODE HERE)
print('last layer output shape: ', last_layer.output_shape)
last_output = last_layer.output ### YOUR CODE HERE
```

```
last layer output shape: (None, 7, 7, 768)
```

```
# Define a Callback class that stops training once accuracy reaches 99.9%
class myCallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        if(logs.get('accuracy')>0.999):
            print("\nReached 99.9% accuracy so cancelling training!")
            self.model.stop_training = True
```

```
from tensorflow.keras.optimizers import RMSprop
```

```
# Flatten the output layer to 1 dimension
x = layers.Flatten()(last_output)
# Add a fully connected layer with 1,024 hidden units and ReLU activation
x = layers.Dense(1024, activation='relu')(x)### YOUR CODE HERE)(x)
# Add a dropout rate of 0.2
x = layers.Dropout(0.2)(x)### YOUR CODE HERE)(x)
# Add a final sigmoid layer for classification
x = layers.Dense (1, activation='sigmoid')(x)### YOUR CODE HERE)(x)
```

```
model = Model( pre_trained_model.input, x) ### YOUR CODE HERE, x)
```

```
model.compile(optimizer = RMSprop(learning_rate=0.0001),
              loss = 'binary_crossentropy', ### YOUR CODE HERE,
              metrics = ['accuracy'] ### YOUR CODE HERE,
              )
```

```
model.summary()
```

activation_155 (Activation) (None, 7, 7, 192) 0 batch_normalization_155[0][0]

activation_155 (Activation)	(None, 7, 7, 192)	0	batch_normalization_155[0][0]
activation_160 (Activation)	(None, 7, 7, 192)	0	batch_normalization_160[0][0]
conv2d_156 (Conv2D)	(None, 7, 7, 192)	258048	activation_155[0][0]
conv2d_161 (Conv2D)	(None, 7, 7, 192)	258048	activation_160[0][0]
batch_normalization_156 (BatchN	(None, 7, 7, 192)	576	conv2d_156[0][0]
batch_normalization_161 (BatchN	(None, 7, 7, 192)	576	conv2d_161[0][0]
activation_156 (Activation)	(None, 7, 7, 192)	0	batch_normalization_156[0][0]
activation_161 (Activation)	(None, 7, 7, 192)	0	batch_normalization_161[0][0]
average_pooling2d_15 (AveragePo	(None, 7, 7, 768)	0	mixed6[0][0]
conv2d_154 (Conv2D)	(None, 7, 7, 192)	147456	mixed6[0][0]
conv2d_157 (Conv2D)	(None, 7, 7, 192)	258048	activation_156[0][0]
conv2d_162 (Conv2D)	(None, 7, 7, 192)	258048	activation_161[0][0]
conv2d_163 (Conv2D)	(None, 7, 7, 192)	147456	average_pooling2d_15[0][0]
batch_normalization_154 (BatchN	(None, 7, 7, 192)	576	conv2d_154[0][0]
batch_normalization_157 (BatchN	(None, 7, 7, 192)	576	conv2d_157[0][0]
batch_normalization_162 (BatchN	(None, 7, 7, 192)	576	conv2d_162[0][0]
batch_normalization_163 (BatchN	(None, 7, 7, 192)	576	conv2d_163[0][0]
activation_154 (Activation)	(None, 7, 7, 192)	0	batch_normalization_154[0][0]
activation_157 (Activation)	(None, 7, 7, 192)	0	batch_normalization_157[0][0]
activation_162 (Activation)	(None, 7, 7, 192)	0	batch_normalization_162[0][0]
activation_163 (Activation)	(None, 7, 7, 192)	0	batch_normalization_163[0][0]
mixed7 (Concatenate)	(None, 7, 7, 768)	0	activation_154[0][0] activation_157[0][0] activation_162[0][0] activation_163[0][0]
flatten (Flatten)	(None, 37632)	0	mixed7[0][0]

dense (Dense)	(None, 1024)	38536192	flatten[0][0]
dropout (Dropout)	(None, 1024)	0	dense[0][0]
dense_1 (Dense)	(None, 1)	1025	dropout[0][0]
=====			
Total params: 47,512,481			
Trainable params: 38,537,217			
Non-trainable params: 8,975,264			
=====			

```
# Get the Horse or Human dataset
!gdown --id 1onaG42NZft3wCE1WH0GDEbUhu75fedP5

# Get the Horse or Human Validation dataset
!gdown --id 1LYeusSEIiZQpwN-mthh5nKdA75VsKG1U

from tensorflow.keras.preprocessing.image import ImageDataGenerator

import os
import zipfile

test_local_zip = './horse-or-human.zip'
zip_ref = zipfile.ZipFile(test_local_zip, 'r')
zip_ref.extractall('./training')

val_local_zip = './validation-horse-or-human.zip'
zip_ref = zipfile.ZipFile(val_local_zip, 'r')
zip_ref.extractall('./validation')

zip_ref.close()

Downloading...
From: https://drive.google.com/uc?id=1onaG42NZft3wCE1WH0GDEbUhu75fedP5
To: /content/horse-or-human.zip
100% 150M/150M [00:01<00:00, 97.2MB/s]
Downloading...
From: https://drive.google.com/uc?id=1LYeusSEIiZQpwN-mthh5nKdA75VsKG1U
To: /content/validation-horse-or-human.zip
100% 11.5M/11.5M [00:00<00:00, 101MB/s]

# Define our example directories and files
train_dir = './training'
```

```

validation_dir = './validation'

# Directory with our training horse pictures
train_horses_dir = os.path.join(train_dir, 'horses') ### YOUR CODE HERE
# Directory with our training humans pictures
train_humans_dir = os.path.join(train_dir, 'humans') ### YOUR CODE HERE
# Directory with our validation horse pictures
validation_horses_dir = os.path.join(validation_dir, 'horses') ### YOUR CODE HERE
# Directory with our validation humanas pictures
validation_humans_dir = os.path.join(validation_dir, 'humans') ### YOUR CODE HERE

train_horses_fnames = os.listdir(train_horses_dir) ### YOUR CODE HERE)
train_humans_fnames = os.listdir(train_humans_dir) ### YOUR CODE HERE)
validation_horses_fnames = os.listdir(validation_horses_dir) ### YOUR CODE HERE)
validation_humans_fnames = os.listdir(validation_humans_dir) ### YOUR CODE HERE)

print(len(train_horses_fnames)) ### YOUR CODE HERE)
print(len(train_humans_fnames)) ### YOUR CODE HERE)
print(len(validation_horses_fnames)) ### YOUR CODE HERE)
print(len(validation_humans_fnames)) ### YOUR CODE HERE)

500
527
128
128

# Add our data-augmentation parameters to ImageDataGenerator
train_datagen = ImageDataGenerator(rescale = 1./255., ### YOUR CODE HERE)
                                rotation_range = 40,
                                width_shift_range = 0.2,
                                height_shift_range = 0.2,
                                shear_range = 0.2,
                                zoom_range = 0.2,
                                horizontal_flip = True)

# Note that the validation data should not be augmented!
test_datagen = ImageDataGenerator( rescale = 1.0/255. ) ### YOUR CODE HERE)

# Flow training images in batches of 20 using train_datagen generator
train_generator = train_datagen.flow_from_directory(train_dir, ### YOUR CODE HERE)
                                batch_size = 20,

```

```

        class_mode = 'binary',
        target_size = (150, 150))

# Flow validation images in batches of 20 using test_datagen generator
validation_generator = test_datagen.flow_from_directory( validation_dir, ### YOUR CODE HERE)
                    batch_size = 20,
                    class_mode = 'binary',
                    target_size = (150, 150))

```

```

Found 1027 images belonging to 2 classes.
Found 256 images belonging to 2 classes.

```

```

# Run this and see how many epochs it should take before the callback
# fires, and stops training at 99.9% accuracy
# (It should take less than 100 epochs)
callbacks = myCallback() ### YOUR CODE HERE)
history = model.fit( ### YOUR CODE HERE)

```

```

    train_generator,
    validation_data = validation_generator,
    steps_per_epoch = 100,
    epochs = 100,
    validation_steps = 5,
    verbose = 2,
    callbacks=callbacks)

```

```
Epoch 1/100
```

```
WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches
100/100 - 13s - loss: 0.0242 - accuracy: 0.9903 - val_loss: 0.0169 - val_accuracy: 0.9900
```



```

import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

```

```
epochs = range(len(acc))
```


```

plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')

```

```
plt.title('Training and validation accuracy')
plt.legend(loc=0)
plt.figure()
```

```
plt.show()
```



<Figure size 432x288 with 0 Axes>

✓ 13s completed at 4:33 PM

